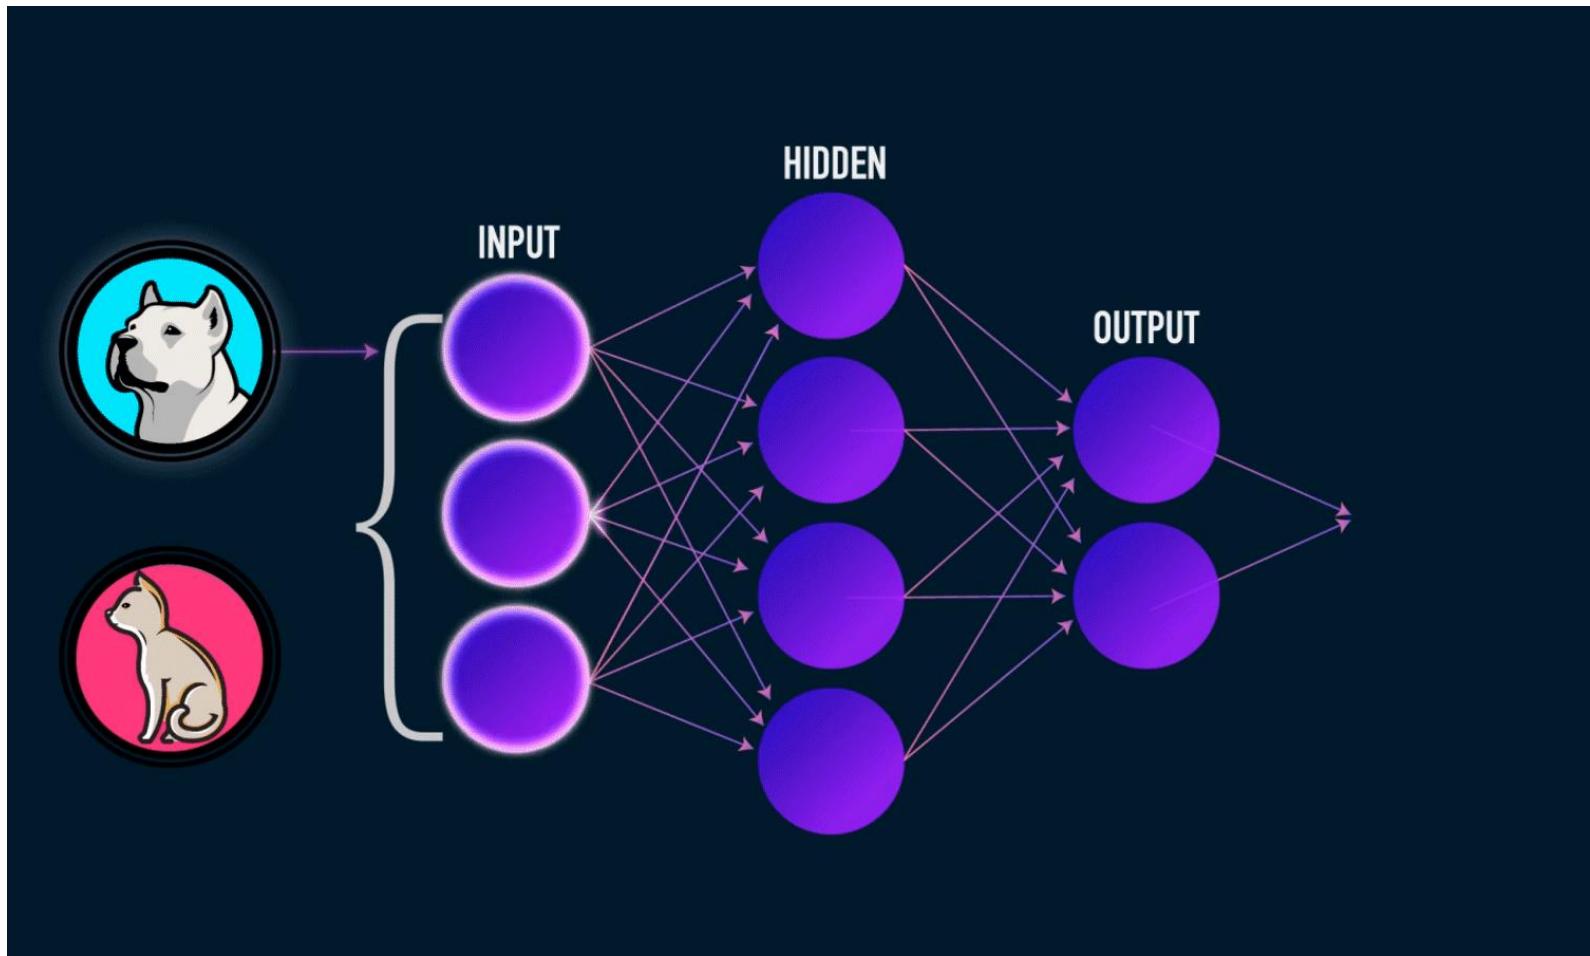


# MLP Math

Simple worked example

# Neural Networks



# What is a Neural Layer?

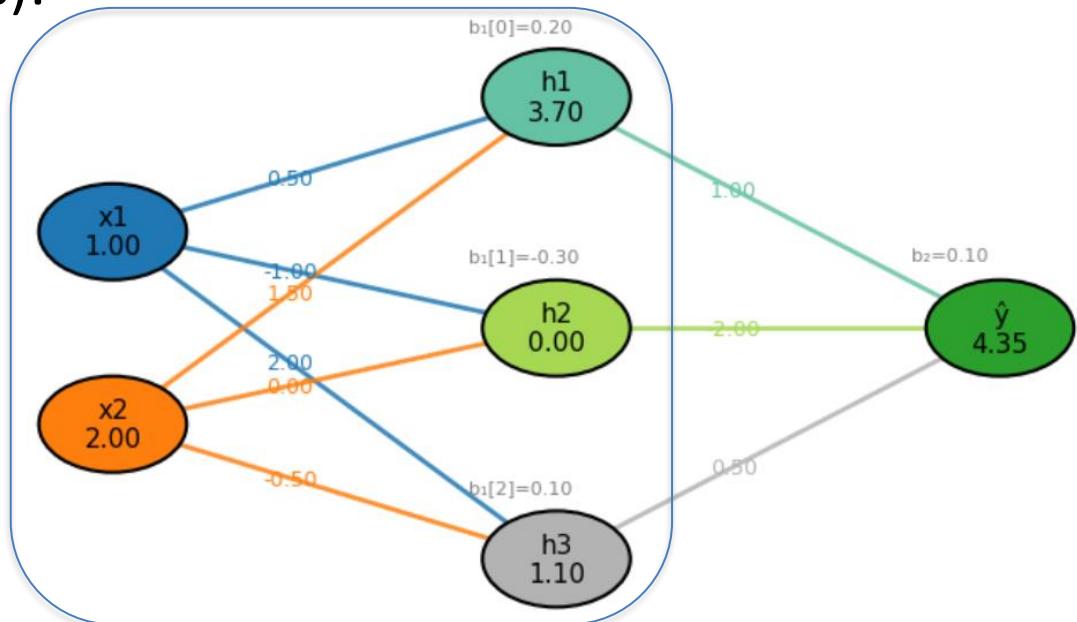
- A layer mixes the input numbers using weights.
- Each output is: weighted sum of inputs + bias.
- Using matrices makes this efficient and compact.

# Matrix Form of a Layer

- Formula:  $Z = XW + b$
- $X$ : inputs ( $N \times d$ )
- $W$ : weights ( $d \times H$ )
- $b$ : bias added to each row ( $1 \times H$ )
- $Z$ : output before activation ( $N \times H$ )

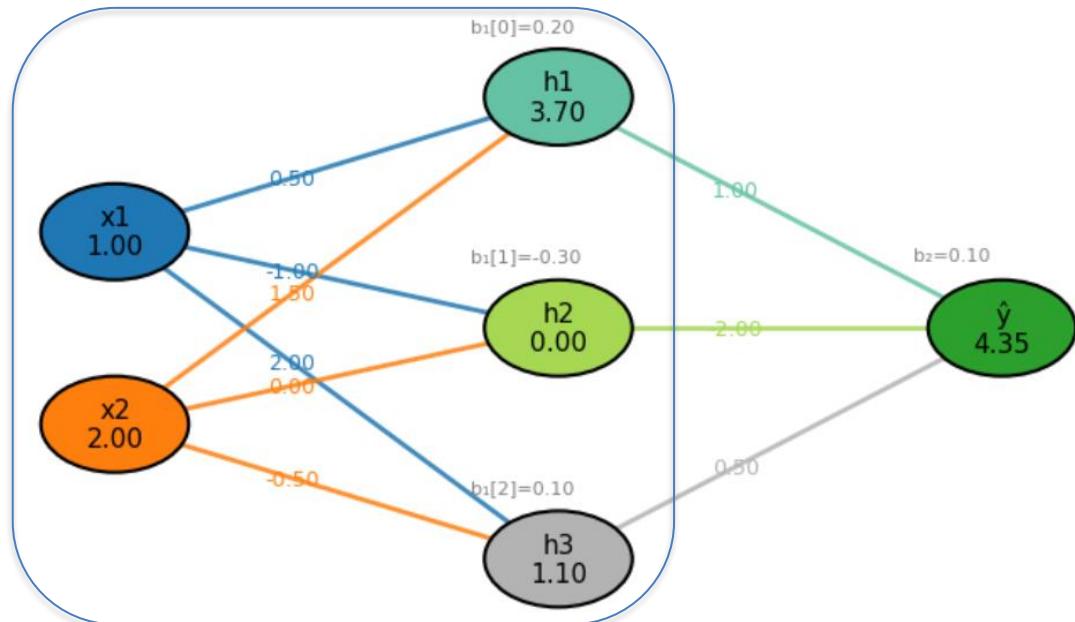
# Example – Input Matrix

- Suppose we have 1 samples with 2 features:
- === INPUT ===
- X (1 sample, 2 features):
- $[[1. 2.]]$



# Example – Weight Matrix

- Let the weight matrix be:
- === LAYER 1 PARAMETERS ===
- W1 (weights from 2 inputs → 3 hidden neurons):
- $\begin{bmatrix} 0.5 & -1. & 2. \end{bmatrix}$
- $\begin{bmatrix} 1.5 & 0. & -0.5 \end{bmatrix}$



# Matrix Multiplication Step

- Compute  $XW$ : (matrix multiplication)

1	2
---	---

Dimension: 1\*2

0.5	-1	2
1.5	0	-0.5

Dimension: 2\*3

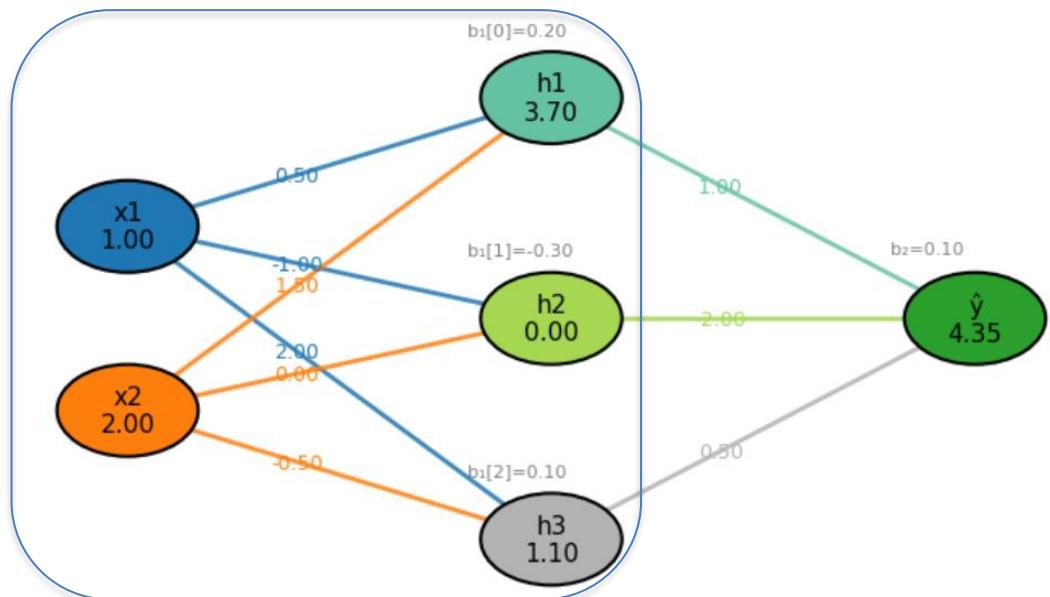
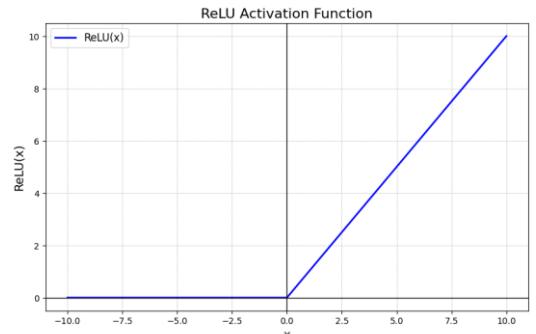
$$XW = \begin{array}{|c|c|c|} \hline 3.5 & -1 & 1 \\ \hline \end{array}$$

# Add Bias

- $b_1$  (bias for each of the 3 hidden neurons):
- $[[ 0.2 \ -0.3 \ 0.1]]$
- $Z = XW + b$  (weighted sums before activation)
- === LAYER 1: LINEAR COMBINATION ===
- $[[ 3.7 \ -1.3 \ 1.1]]$

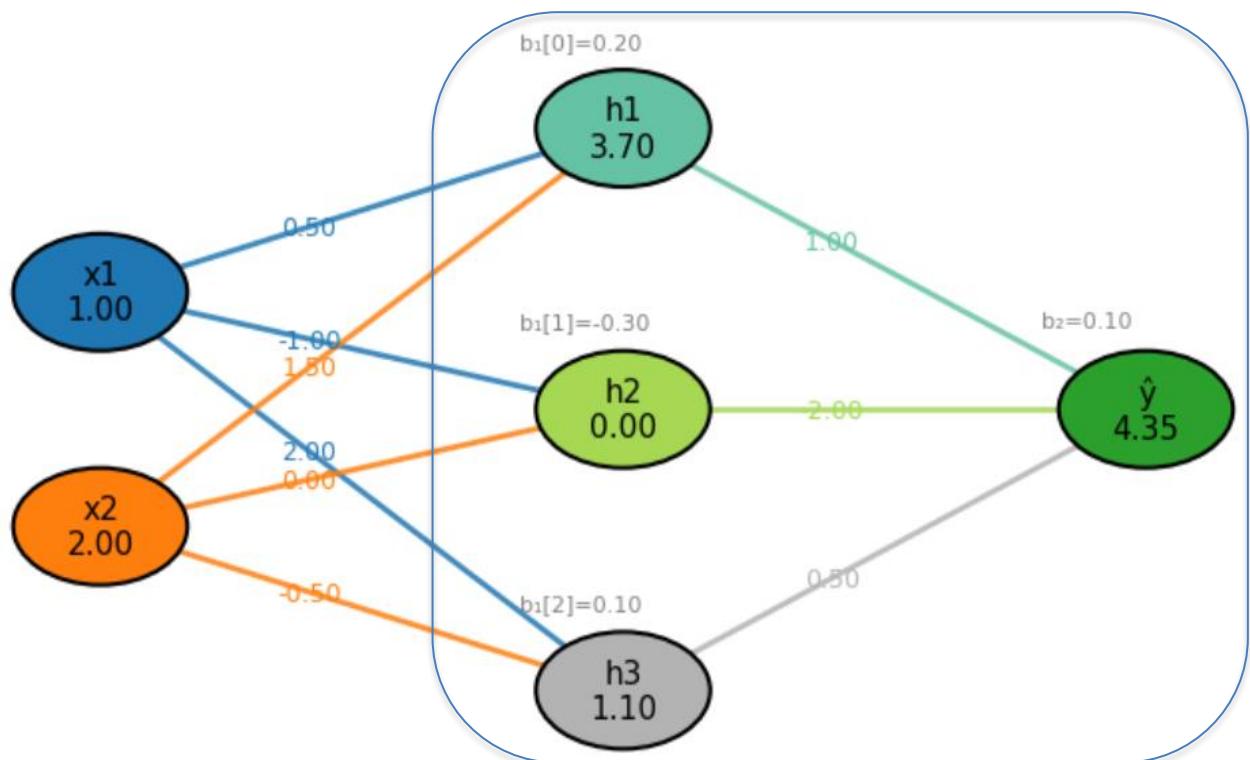
# Activation: ReLU

- $\text{ReLU}(x) = \max(0, x)$
- This produces the hidden layer output.
- === LAYER 1: ACTIVATION ===
- $H1 = \text{ReLU}(Z1)$
- $[[3.7 \ 0. \ 1.1]]$



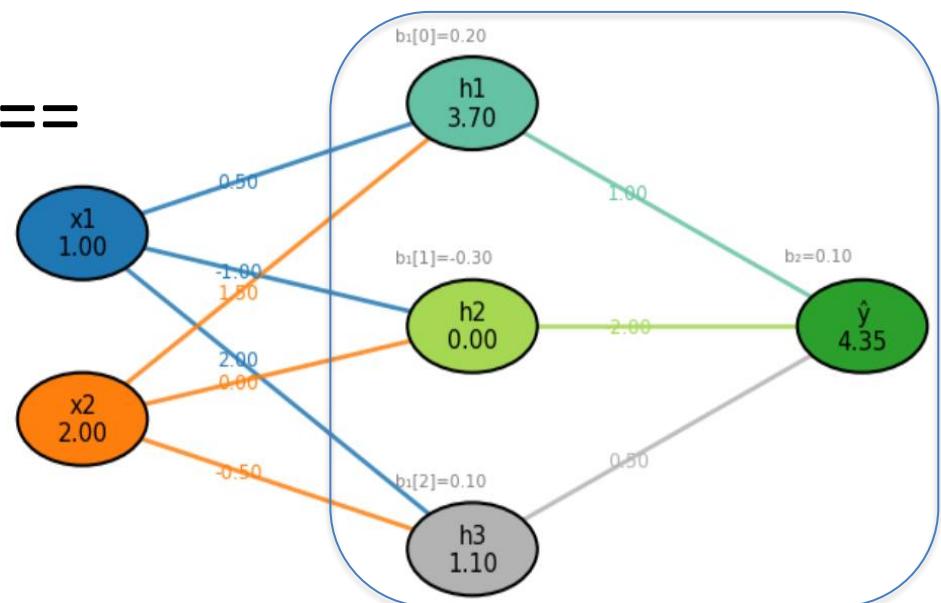
# Layer 2

- $W_2$  (weights from 3 hidden neurons  $\rightarrow$  1 output):
- $[[ 1. ] [-2. ] [ 0.5]]$
- $b_2$  (output bias):
- $[[0.1]]$



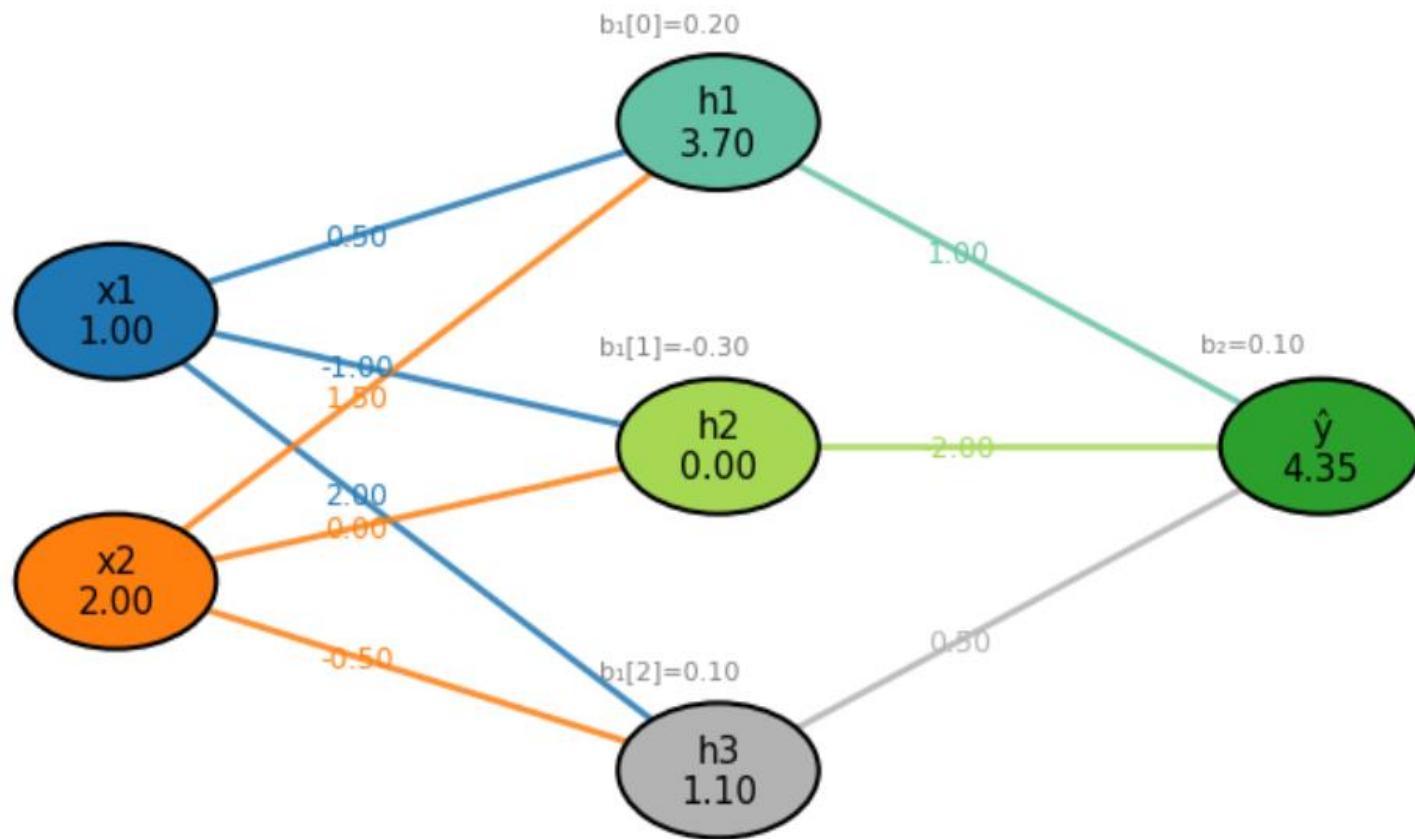
# Layer 2 activation function

- Linear activation function  $f(x) = x$
- === OUTPUT LAYER ===
- $Z_2 = H_1 \cdot W_2 + b_2$  (final prediction before activation):  
[[4.35]]
- === FINAL OUTPUT ===
- Predicted y value:  
[[4.35]]



# Final Regression Result

- $X = [[1, 2]] \rightarrow \hat{y} = 4.35$



# Backpropagation

Step-by-step

# What is Backpropagation?

- A method to find how much each weight contributed to the error.
- Goal: reduce the error by adjusting weights.
- Flow: output → error → send error backward.

# Step 1 – Forward Pass

- Inputs move through the network.
- Each neuron computes:  $(\text{inputs} \times \text{weights}) + \text{bias}$ .
- Activation function is applied.
- We get the prediction  $\hat{y}$ .

# Step 2 – Compute the Error (Loss)

- Compare network output ( $\hat{y}$ ) with target ( $y$ ).
- Loss = how wrong the network is.
- Example:  $MSE = (\hat{y} - y)^2$ .

# Step 3 – Delta at Output Layer

- Delta = how much the output neuron caused the error.
- Comes from derivative of the loss.
- This signal goes backward.

# Step 4 – Gradients for Output Weights

- Core rule: gradient = (error delta) × (input value).
- Large neuron value → large influence → larger update.
- Small or zero neuron value → small influence.

# Step 5 – Send Error Back to Hidden Layer

- Output delta is multiplied by connecting weights.
- Each hidden neuron receives a portion of the error.
- This becomes its delta.

# Step 6 – Activation Derivative

- Hidden neurons apply derivative of activation function.
- ReLU derivative: 1 for positive input, 0 for negative.
- If derivative = 0 → neuron gets no error → no update.

# Step 7 – Gradients for First-Layer Weights

- Same rule: gradient = input  $\times$  delta of hidden neuron.
- Shows how much each weight contributed to final error.

# Step 8 – Bias Gradients

- Bias has constant input = 1.
- So gradient of bias = delta of that neuron.

# Step 9 – Weight Update

- Weights change to reduce error:
- $\text{new\_weight} = \text{old\_weight} - \text{learning\_rate} \times \text{gradient}$ .
- Repeating this trains the network.

# Backprop Summary

- Forward: compute outputs.
- Loss: compare with target.
- Backward: compute deltas.
- Gradients = input  $\times$  delta.
- Update weights to reduce error.
- This is how neural networks learn.

# Forward pass

## 1) Forward pass (with one sample)

Given

$$X = [1, 2]$$

$$W_1 = \begin{bmatrix} 0.5 & -1 & 2 \\ 1.5 & 0 & -0.5 \end{bmatrix}, \quad b_1 = [0.2, -0.3, 0.1]$$

$$W_2 = \begin{bmatrix} 1 \\ -2 \\ 0.5 \end{bmatrix}, \quad b_2 = [0.1]$$

$$y = 3.5$$

Hidden pre-activation  $z_1 = XW_1 + b_1$

- First hidden:  $1 \cdot 0.5 + 2 \cdot 1.5 + 0.2 = 0.5 + 3 + 0.2 = 3.7$
- Second hidden:  $1 \cdot (-1) + 2 \cdot 0 + (-0.3) = -1 - 0.3 = -1.3$
- Third hidden:  $1 \cdot 2 + 2 \cdot (-0.5) + 0.1 = 2 - 1 + 0.1 = 1.1$

So  $z_1 = [3.7, -1.3, 1.1]$

Hidden activation  $h_1 = \text{ReLU}(z_1)$

$$\Rightarrow h_1 = [3.7, 0, 1.1]$$

Output  $z_2 = h_1^\top W_2 + b_2$

$$= 3.7 \cdot 1 + 0 \cdot (-2) + 1.1 \cdot 0.5 + 0.1 = 3.7 + 0 + 0.55 + 0.1 = 4.35$$

Prediction  $\hat{y} = 4.35$

# Backward pass - error

MSE = mean of squared errors

## 2) Loss (MSE)

With one sample,  $L = (\hat{y} - y)^2 = (4.35 - 3.5)^2 = 0.85^2 = \boxed{0.7225}$

# Backward pass – gradient w2

## 3) Output gradient

For MSE with  $\hat{y} = z_2$ :

$$\frac{\partial L}{\partial \hat{y}} = 2(\hat{y} - y) = 2 \cdot 0.85 = \boxed{1.7}$$

Since  $\hat{y} = z_2$ ,  $\frac{\partial L}{\partial z_2} = \boxed{1.7}$ .

Bias gradient at output:  $\boxed{db_2 = 1.7}$

Weights to output:

$$dW_2 = h_1^\top \cdot \frac{\partial L}{\partial z_2} \text{ (outer product with a scalar)}$$

$$dW_2 = \begin{bmatrix} 3.7 \\ 0 \\ 1.1 \end{bmatrix} \cdot 1.7 = \boxed{\begin{bmatrix} 6.29 \\ 0 \\ 1.87 \end{bmatrix}}$$

# Backpropagation to hidden

## 4) Backprop to hidden

First, push error back through  $W_2$ :

$$dH_1 = \frac{\partial L}{\partial z_2} W_2^\top = 1.7 [1, -2, 0.5] = [1.7, -3.4, 0.85]$$

Apply ReLU derivative (1 if  $z_1 > 0$ , else 0):

$$z_1 = [3.7, -1.3, 1.1] \Rightarrow \text{ReLU}' = [1, 0, 1]$$

Hidden deltas:

$$\delta_1 = dZ_1 = dH_1 \odot \text{ReLU}' = [1.7, 0, 0.85]$$

Bias gradients at hidden (mean over batch of size 1):

$$db_1 = [1.7, 0, 0.85]$$

# Backward pass – gradient w1

## 5) Gradients for $W_1$

$$dW_1 = X^\top \cdot \delta_1 \text{ (with } X = [1, 2])$$

Row 1 (from input  $x_1 = 1$ ):  $1 \cdot [1.7, 0, 0.85] = [1.7, 0, 0.85]$

Row 2 (from input  $x_2 = 2$ ):  $2 \cdot [1.7, 0, 0.85] = [3.4, 0, 1.7]$

$$dW_1 = \begin{bmatrix} 1.7 & 0 & 0.85 \\ 3.4 & 0 & 1.7 \end{bmatrix}$$

# Final Results

## 6) Final collected results

- Loss (MSE):  $0.7225$
- $dW_1:$   $\begin{bmatrix} 1.7 & 0 & 0.85 \\ 3.4 & 0 & 1.7 \end{bmatrix}$
- $dW_2:$   $\begin{bmatrix} 6.29 \\ 0 \\ 1.87 \end{bmatrix}$
- $db_1:$   $[1.7, 0, 0.85]$
- $db_2:$   $[1.7]$

# Why Matrix Multiplication?

- Processes many samples at once.
- More efficient than doing each weighted sum manually.
- Matches how real neural networks run on GPUs/CPUs.